

# Behavioral Cloning for Lunar Lander

Naya Baslan, Megan Klaiber, Ahmed Krimi, Abdelrahman Younes

## Project Description

Lunar Lander is an OpenAIGym Environment whose goal is to navigate a space-ship and make it land between the two yellow flags with up right position and minimum landing velocity. As shown in the figure, there are four discrete action classes: upwards, left, right and no-action.

In this project, we approach this task through imitation learning, where a human plays the game and feeds the training data so that the agent can learn from it, thereby reformulating the problem as a supervised learning problem.

## Obtaining the Training Data

- This was done by collecting data from a human playing the lunar lander game
- Total samples collected: 27500

## Training the Behavioral Cloning Agent

- Experiments were performed by implementing a fully connected network (FCN) and a convolutional neural network (CNN).
- FCN takes the state (e.g. angle, velocity) as input whereas the CNN learns directly from images.




Figure 1: Lunar Lander Environment

## Fully Connected Neural Network

### Network architecture:

- Input Layer: 8 states  
[ x\_pos, y\_pos, x\_vel, y\_vel, angle, ang\_vel, left\_leg\_ground, right\_leg\_ground ]
- 3 fully connected layers with dropout (0.5) and activation functions ReLU, cross-entropy loss
- Output: one of four action classes (right, left, up, no action)
- Optimizer: Adam with learning rate = 0.001

Table 1: Test result over 15 test episodes

	mean	std
FCN	26.83	22.34

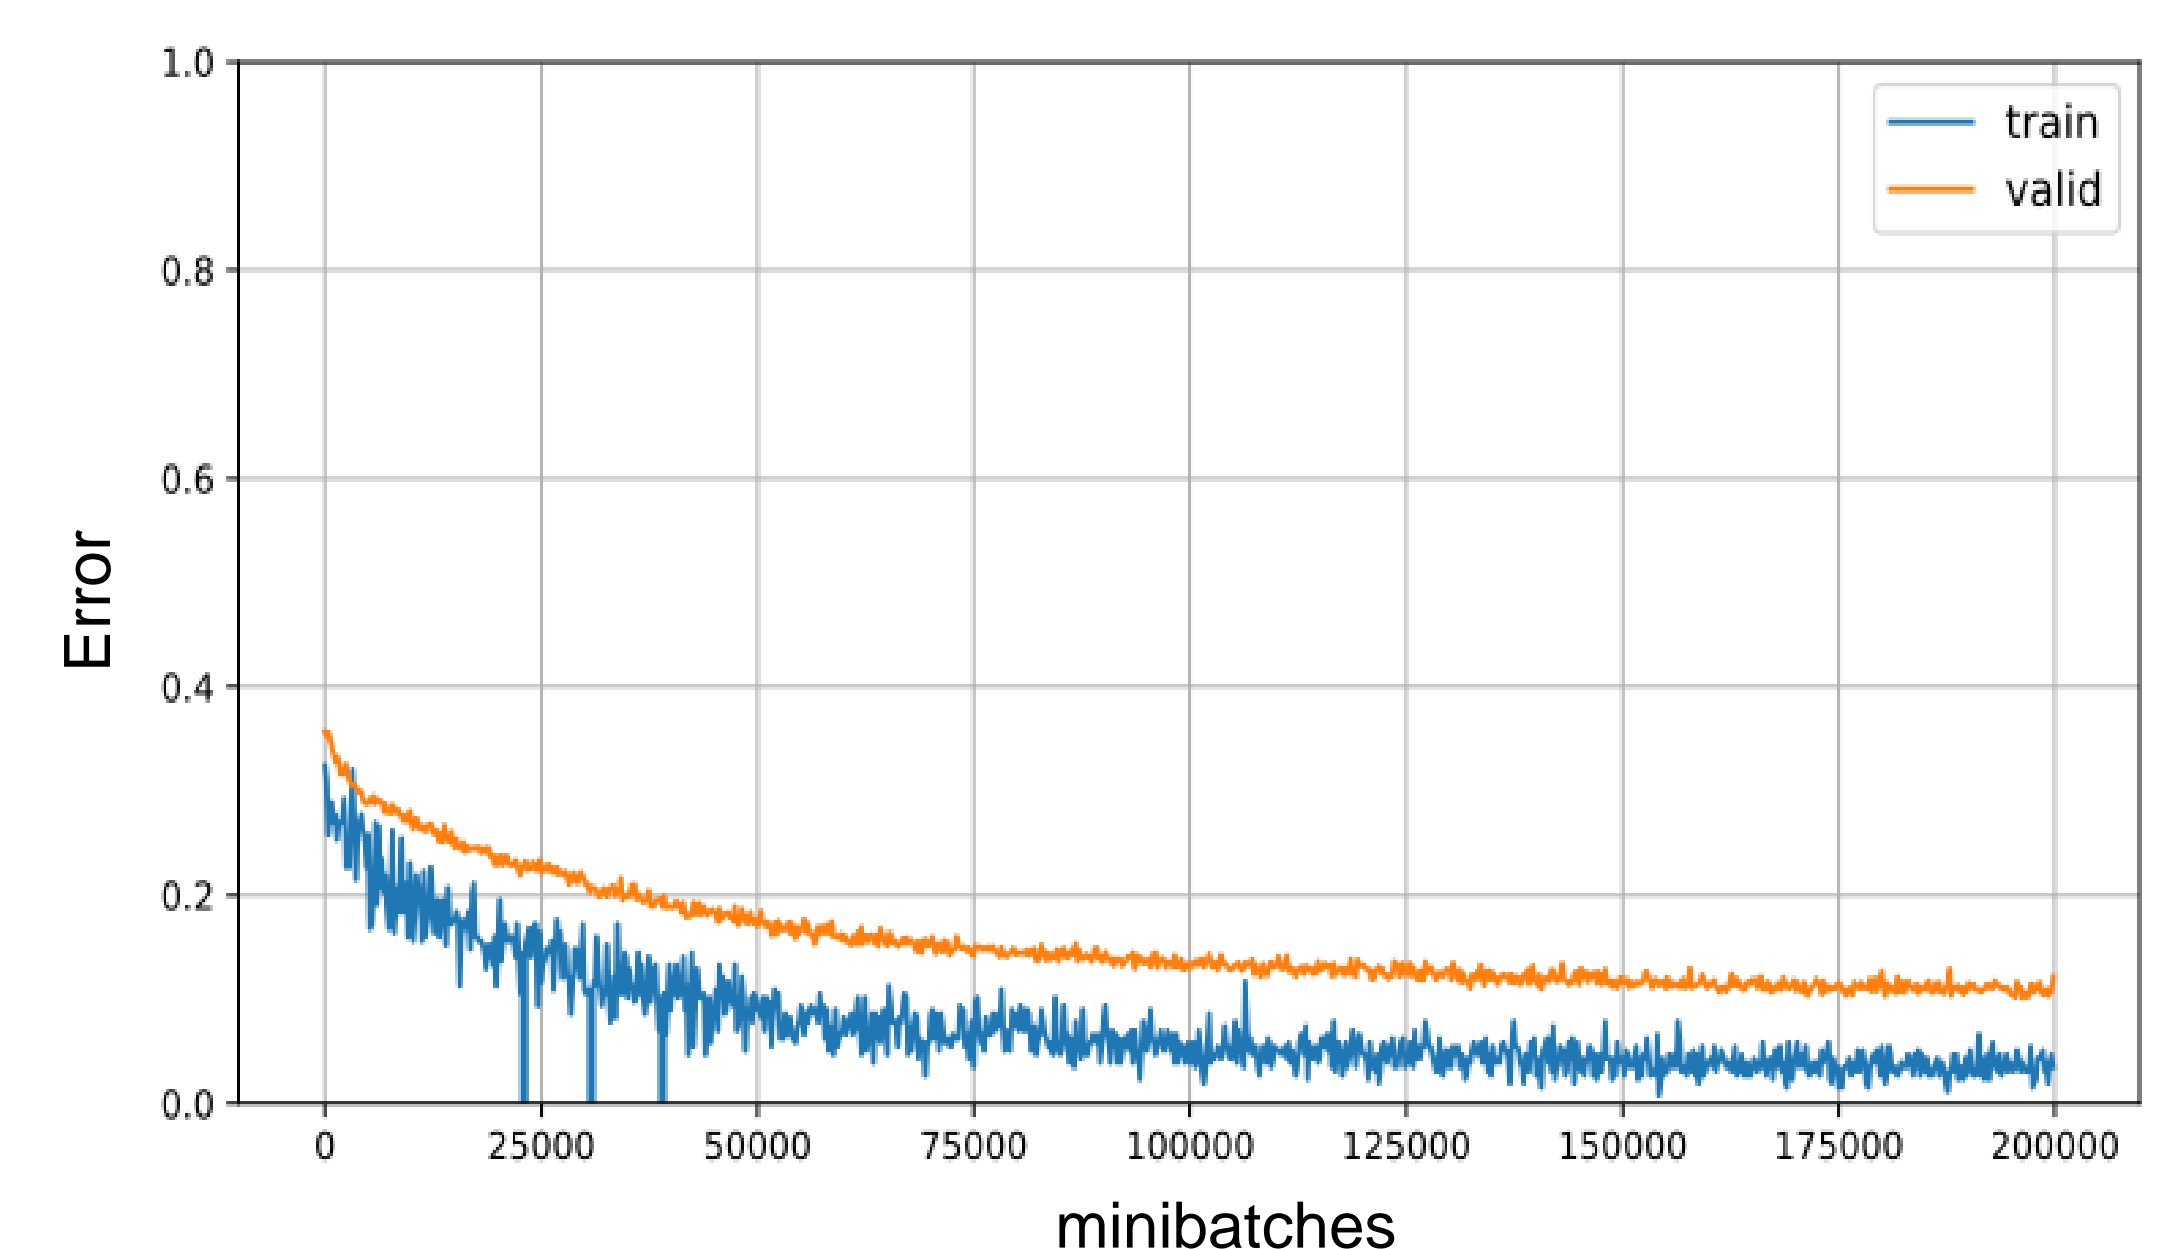


Figure 2: Training and Validation of FCN

## Convolutional Neural Network (5 different architectures)

### 1. First approach: CNN with 2 convolutional layers

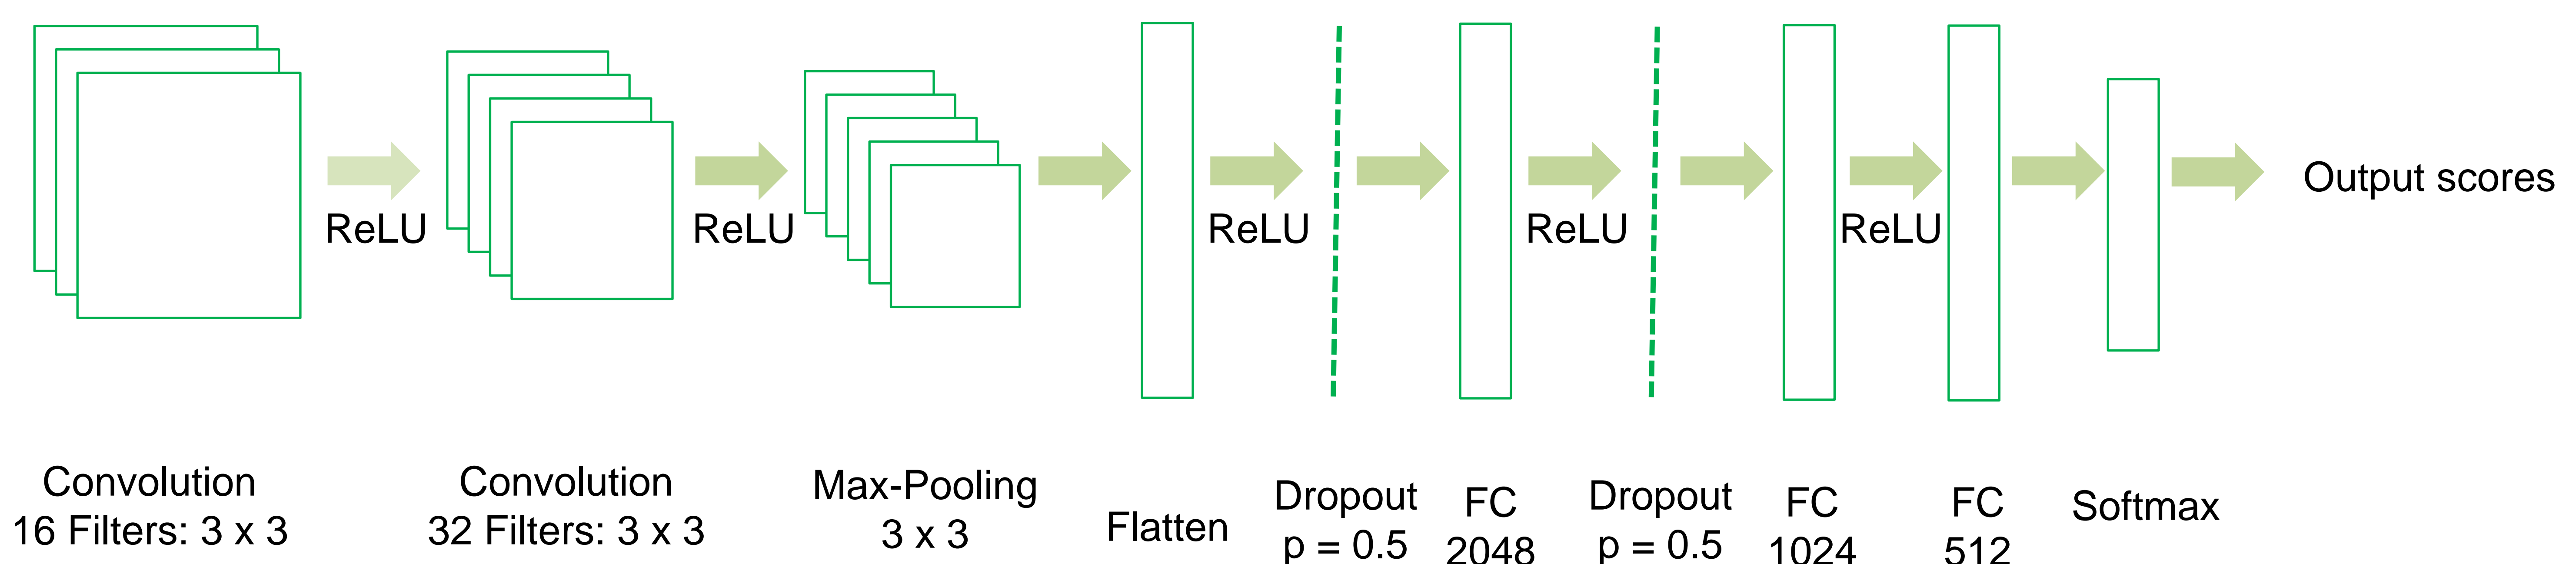


Figure 3: CNN Architecture

### Other hyperparameter settings:

Learning Rate = 0.001, Batch Size = 64, Optimizer = Adam, Activation Function = ReLU, #Minibatches: 1400, Stride = 1, Loss = Cross Entropy Loss



# Behavioral Cloning for Lunar Lander

## 2. Second approach: CNN with 5 convolutional layers + uniform sampling

The collected data is highly imbalanced because the no-action class occurs more frequently than the others. Uniform sampling was implemented here to equally sample between the different action classes.

### Network Architecture Description:

5 convolutional layers were used with respective filter sizes =  $[7 \times 7]$ ,  $[5 \times 5]$ ,  $[5 \times 5]$ ,  $[3 \times 3]$ ,  $[3 \times 3]$ , followed by 3 fully-connected layers.

### Other Hyperparameter Settings:

Learning Rate = 0.001, Batch Size = 128, Optimizer = Adam, Activation Functions = ReLU,  
#Minibatches: 1500, Stride = 1, Dropout = 0.8, Loss = Cross Entropy Loss, Number of Filters = 32

## 3. Third Approach: Modified ZFNet: (Different Configurations - simplified to 6 layers)

Different Sampling Methods were used here:

- Equal Sampling, No-action limiter to one fourth and one half of the training set size

### Other Hyperparameter Settings:

Learning Rate = 0.001, Batch Size = 128, Optimizer = Adam  
Activation Functions = ReLU, #Minibatches: 1500, Stride = 1, Dropout = 0.5, Loss = Cross Entropy  
Loss, Kernel size =  $[7 \times 7]$ ,  $[3 \times 3]$

## 4. Fourth Approach: Modified ResNet18:

### Network Architecture Description:

Similar to ResNet18 architecture; remove average pooling, and reduce the number of hidden layers to 4

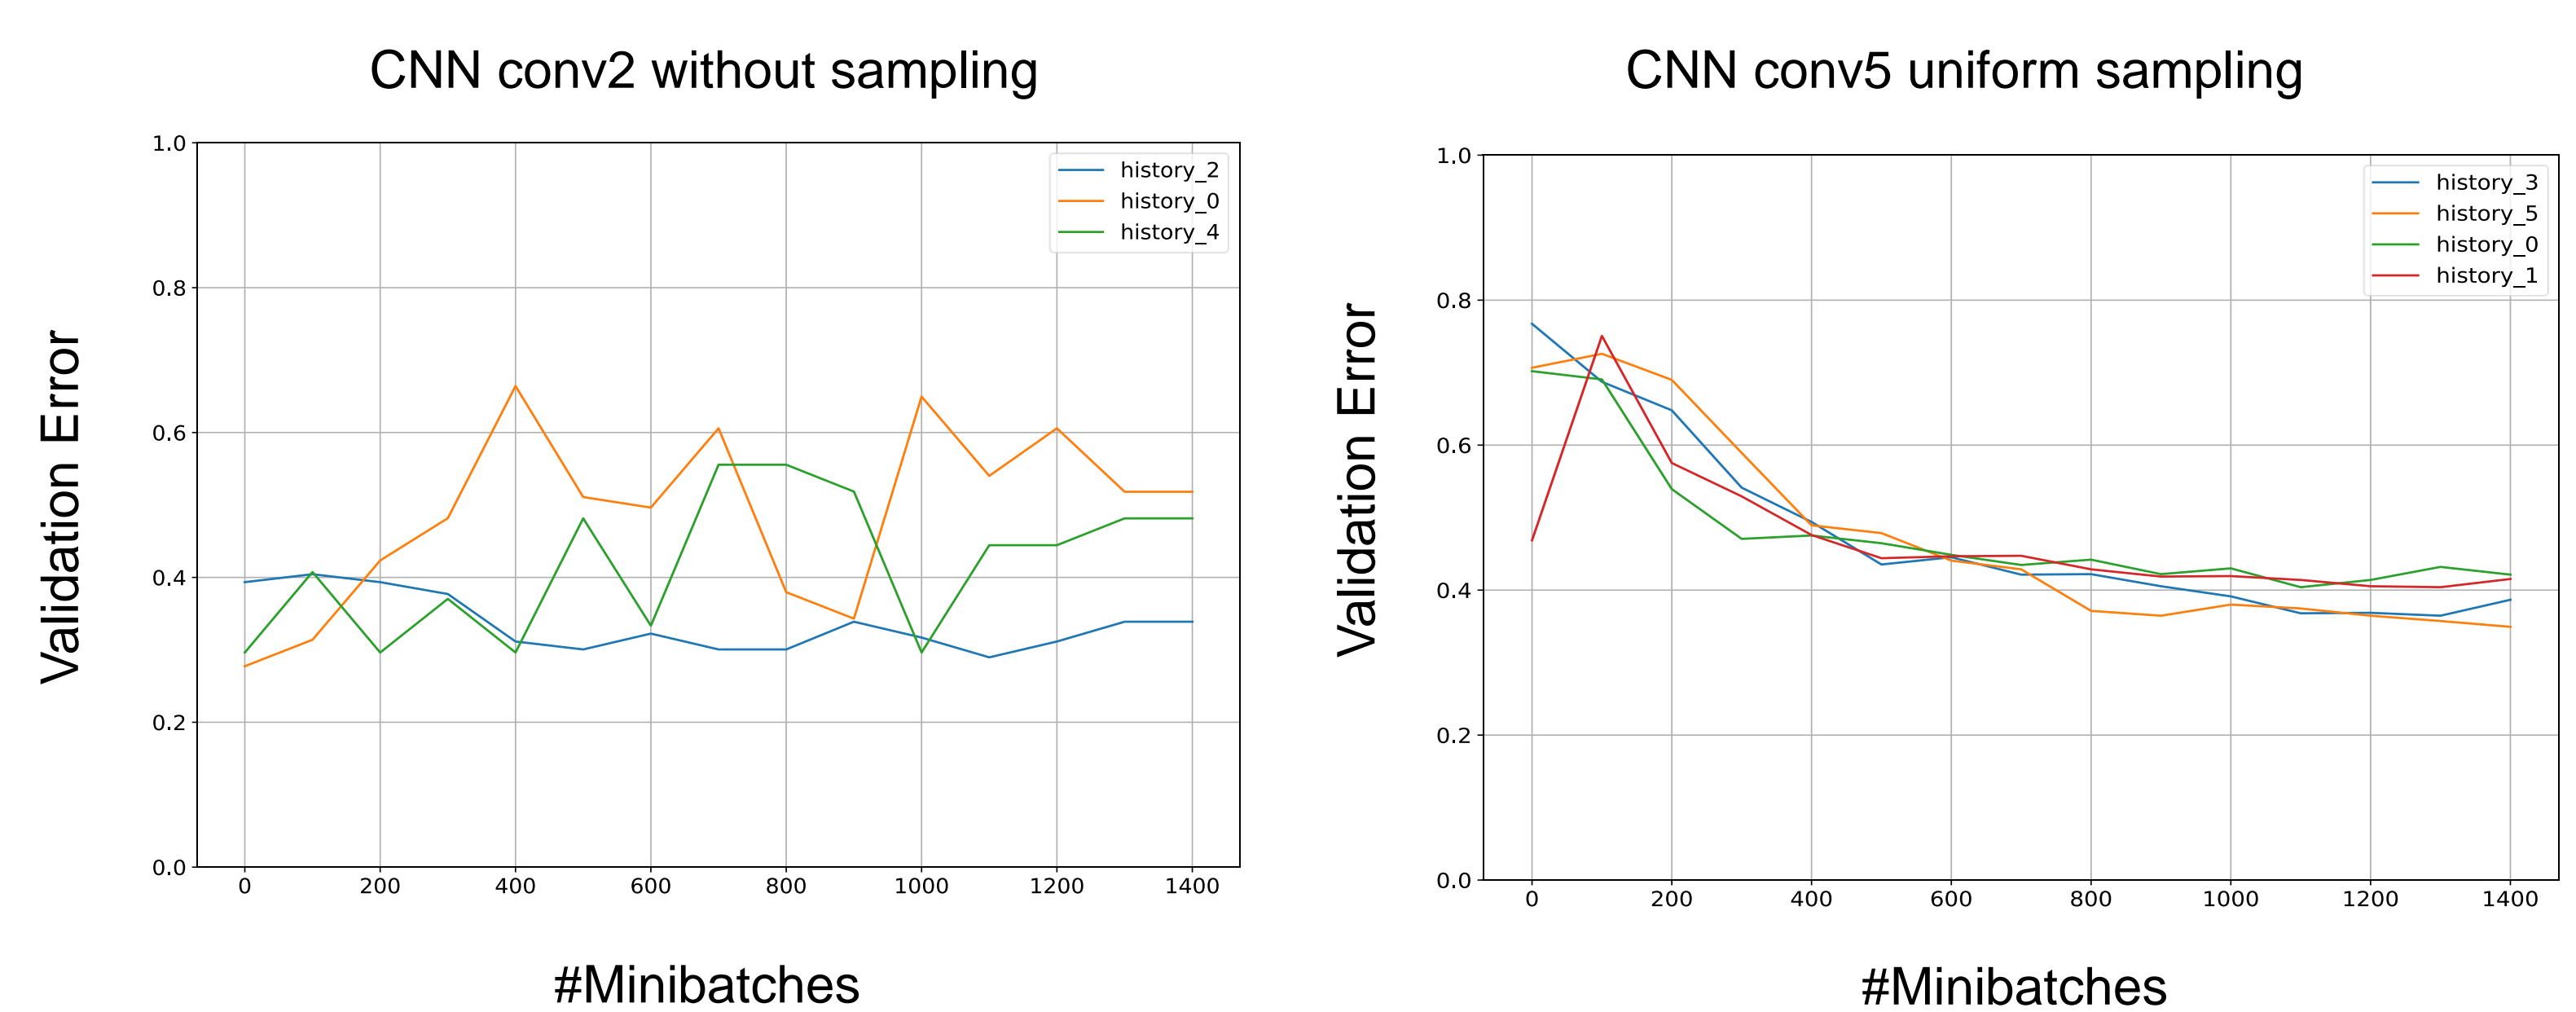
### Other Hyperparameter Settings:

Learning Rate = 0.0005, Batch Size = 16, Optimizer = Adam, Activation Functions = ReLU,  
#Minibatches: 1200, Stride = 1, Loss = Cross Entropy Loss

## Experimental Results for the Convolutional Neural Networks

**Table 2: Mean performance and standard deviation over 15 test episodes**

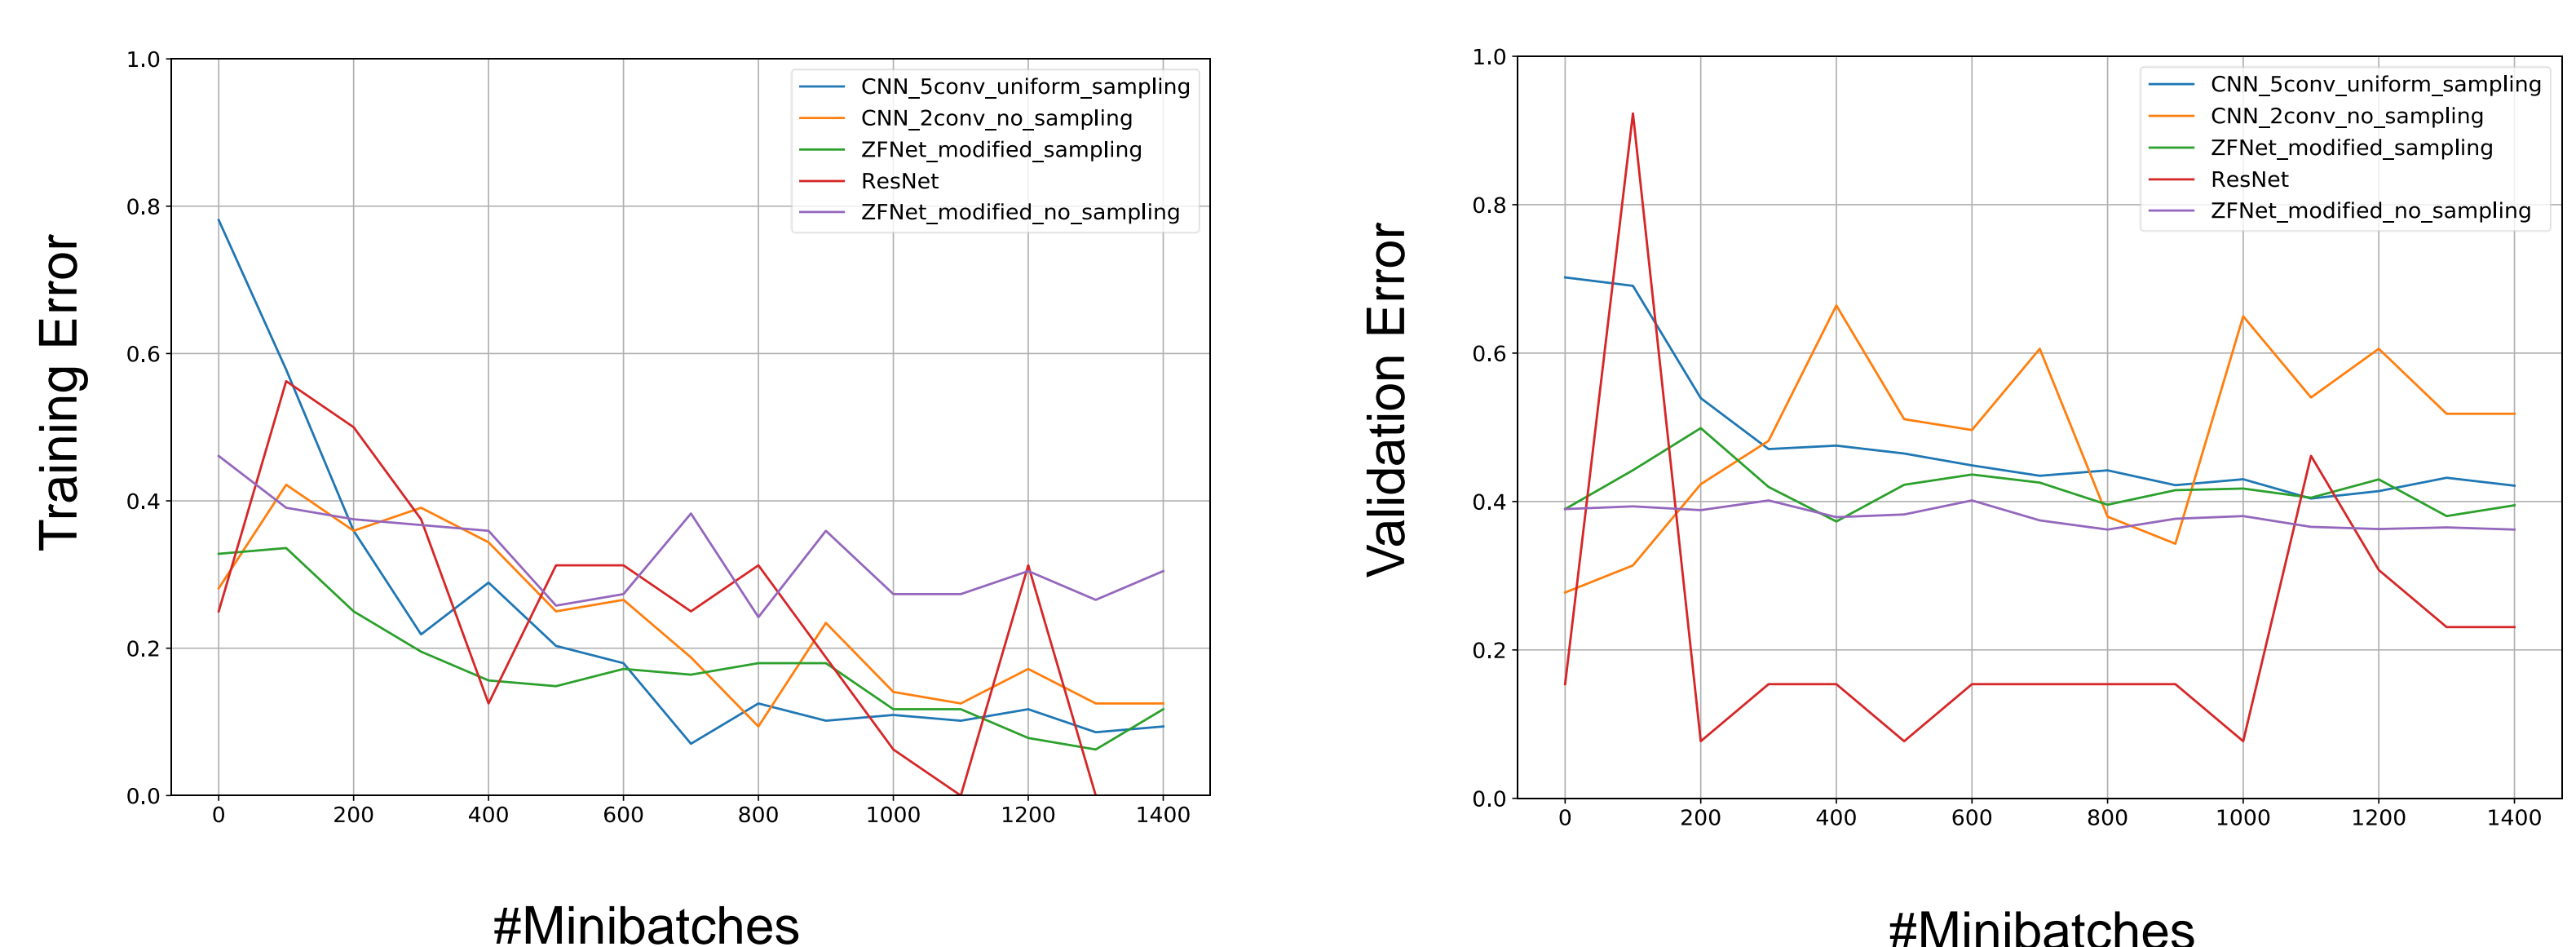
CNN 2 convs	mean	std	CNN 5 convs	mean	std
history length 0	12.76	26.88	history length 0	14.34	20.80
history length 2	28.33	24.38	history length 1	-11.34	2.21
history length 4	4.56	28.10	history length 3	8.14	22.35
			history length 5	8.70	18.50



**Figure 4: Validation of two CNN Architectures using History**

## Conclusion

- FCN performs quite well in testing but requires more time to train
- CNN has the best performance for history  $N = 2$
- LSTMs were also implemented and seem to have preliminary promising results
- Attached video shows the performance of the different agents



**Figure 5: Training and Validation of all five CNN Architectures without history**